

Characterizing and Evaluating Adversarial Examples in Handwritten Signature Verification

Team Members:

Challapalli Sathwik([Sathwikchallapalli](#))

Talasila Revanth([revanth7375](#))

B Sanjeev Roy([sanjeevroy08](#))

Problem Statement

With the increasing use of digital documentation and automated verification systems, ensuring the authenticity of handwritten signatures has become critically important. Traditional manual verification methods are prone to human error, time-consuming, and not scalable to large datasets. Machine learning and deep learning models provide automated solutions for signature verification but remain vulnerable to adversarial attacks—small, intentional changes or perturbations to input images that can cause misclassification.

This project addresses the problem of adversarial vulnerability in deep learning-based signature verification systems. Using a ResNet50-based convolutional neural network, the goal is to accurately classify genuine and forged signatures while assessing and improving the model's resilience against adversarial attacks such as the Fast Gradient Sign Method (FGSM). The system aims to:

1. Detect forged signatures with high accuracy.
2. Analyze how adversarial noise impacts model predictions.
3. Implement adversarial training to make the network robust against such attacks.

The final outcome is a robust and reliable signature verification model capable of maintaining performance even when exposed to manipulated adversarial inputs, thereby strengthening the security of biometric authentication systems.

Objectives

- Develop a CNN-based signature verification model using pretrained ResNet50.
- Preprocess signature images for optimal input (grayscale, resize, binarize, normalize).
- Generate adversarial examples using the Fast Gradient Sign Method (FGSM).
- Evaluate model vulnerability to adversarial attacks.
- Improve model robustness via adversarial training with clean and perturbed data.
- Measure and compare performance pre- and post-robustification using accuracy and ROC-AUC.
- Visualize adversarial perturbations and their effect on predictions.

Dataset used

The CEDAR Signature Dataset is a widely-used benchmark dataset for handwritten signature verification. Here are some key details:

- Contains signatures from 55 individuals (writers).
- Each individual provides 24 genuine signature samples and 24 forged signatures.
- The total dataset has 2,640 signature images.
- The signatures were scanned at 300 dpi in grayscale and binarized.
- Preprocessing usually involves noise removal and normalization to improve model accuracy.
- This dataset is ideal for offline signature verification (working on scanned images rather than online stroke data).

It is commonly used in academic research as a standard for testing signature verification algorithms, such as CNN-based models and adversarial robustness methods

Solution

The solution proposed in this project is a deep learning-based signature verification system that utilizes the pretrained ResNet50 convolutional neural network architecture to accurately classify genuine and forged handwritten signatures. The model is trained on a curated and preprocessed signature dataset, including grayscale conversion, binarization, resizing, and normalization. To enhance security, the system incorporates adversarial testing using the Fast Gradient Sign Method (FGSM) to assess vulnerabilities to malicious input perturbations. Subsequently, adversarial training is employed by exposing the model to both clean and adversarial examples, thereby increasing its robustness against such attacks. The final system demonstrates high accuracy and resilience, making it suitable for deployment in biometric authentication applications where reliable signature verification is crucial.

Algorithm

METHODOLOGY FOR CHARACTERIZING AND EVALUATING ADVERSARIAL EXAMPLES IN HANDWRITTEN SIGNATURE VERIFICATION

1. Dataset Collection:
Download signature datasets such as CEDAR or Kaggle Signature Dataset containing genuine and forged signature images.
2. Preprocessing:
 - Convert images to grayscale.
 - Resize to 224×224 pixels for ResNet50 input.
 - Apply OTSU thresholding to binarize signatures.
 - Convert grayscale images to 3-channel RGB format.
 - Normalize images using ResNet50's preprocess_input function.
3. Model Construction:
 - Load ResNet50 pretrained on ImageNet without the top classification layer.
 - Add Global Average Pooling layer to convert feature maps to vector.
 - Add Dropout with 0.5 rate to mitigate overfitting.
 - Add output Dense layer with sigmoid activation for binary classification.
4. Training:
 - Freeze the ResNet50 base layers initially.
 - Train the newly added top layers on clean training dataset.
 - Fine-tune the last 20 layers of ResNet50 with a smaller learning rate for improved performance.
5. Adversarial Attack (FGSM):
 - Implement FGSM to generate adversarial samples by perturbing input images along the gradient of the loss function.
 - Test the trained model on these adversarial examples to evaluate its vulnerability.

6. Adversarial Training:

- Retrain the model by mixing clean and adversarial samples in each batch.
- Update weights to improve robustness to adversarial perturbations.

7. Evaluation:

- Evaluate model performance on clean and adversarial test sets using accuracy, precision, recall, F1-score, and ROC-AUC metrics.
- Visualize original and adversarial images to illustrate subtle perturbations affecting model predictions.

8. Saving Model:

- Save the robust adversarially trained model weights for future use or deployment.

ARCHITECTURE DIAGRAM FOR RESNET-50

ResNet-50 Architecture



PROGRAMS & OUTPUTS:

MODEL BUILDING FOR SIGNATURE VERIFICATION CODE:

```
import zipfile
import os
import numpy as np
import cv2
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.applications.resnet50 import ResNet50,
preprocess_input
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, roc_auc_score
from tensorflow.keras.callbacks import ModelCheckpoint
```

```

# 1. UNZIP YOUR UPLOADED DATASET
zip_file_path = '/content/signatures.zip' # Change if different!
extract_dir = 'SignatureDataset'

print("Extracting zip file...")
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(extract_dir)

dataset_dir = 'SignatureDataset/signatures'

print("Extracted folders:", os.listdir(dataset_dir))
print("Genuine examples:", os.listdir(os.path.join(dataset_dir,
'Genuine'))[:3])
print("Forgery examples:", os.listdir(os.path.join(dataset_dir,
'Forgery'))[:3])

# 2. IMAGE PREPROCESSING FUNCTION
def preprocess_signature_image(image_path):
    image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
    image = cv2.resize(image, (224, 224)) # For ResNet50 input size
    _, thresh_img = cv2.threshold(image, 0, 255, cv2.THRESH_BINARY +
cv2.THRESH_OTSU)
    inverted = cv2.bitwise_not(thresh_img)
    # Convert to 3 channels
    img_3ch = cv2.cvtColor(inverted, cv2.COLOR_GRAY2RGB)
    return img_3ch

# 3. LOAD DATASET IMAGES AND LABELS
def load_dataset(dataset_dir):
    X = []
    y = []
    for label_name, label_value in [('Genuine', 1), ('Forgery', 0)]:
        folder_path = os.path.join(dataset_dir, label_name)
        if not os.path.exists(folder_path):
            print(f"Warning: folder not found: {folder_path}")
            continue
        for filename in os.listdir(folder_path):
            if filename.lower().endswith('.png', '.jpg', '.jpeg',
'.bmp')):
                img_path = os.path.join(folder_path, filename)
                img = preprocess_signature_image(img_path)
                X.append(img)
                y.append(label_value)
    return np.array(X), np.array(y)

# 4. LOAD AND PREPROCESS DATA
print("\nLoading and preprocessing dataset...")
X, y = load_dataset(dataset_dir)

```

```
print(f"Loaded {len(X)} samples.")

# Normalize using preprocess_input from ResNet50
X = preprocess_input(X.astype(np.float32))

# Split dataset
print("\nSplitting into train/test...")
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42)

# 5. DEFINE CNN MODEL BASED ON ResNet50
base_model = ResNet50(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
predictions = Dense(1, activation='sigmoid')(x)
model = Model(inputs=base_model.input, outputs=predictions)

# Freeze base model layers first
for layer in base_model.layers:
    layer.trainable = False

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

# Define checkpoint callback to save best model during initial training
# (new Keras format)
checkpoint_path = 'best_model_initial.keras'
checkpoint = ModelCheckpoint(checkpoint_path, monitor='val_accuracy',
verbose=1,
                           save_best_only=True, mode='max')

print("\nTraining CNN model...")
history = model.fit(
    X_train, y_train,
    validation_split=0.2,
    epochs=5,
    batch_size=32,
    verbose=2,
    callbacks=[checkpoint]
)

# Load the best weights from initial training
model.load_weights(checkpoint_path)

# Optionally, unfreeze some layers and fine-tune further
for layer in base_model.layers[-20:]:
```

```

        layer.trainable = True
model.compile(optimizer=tf.keras.optimizers.Adam(1e-5),
loss='binary_crossentropy', metrics=['accuracy'])

# Define checkpoint callback for fine-tuning phase (new Keras format)
fine_tune_checkpoint_path = 'best_model_finetuned.keras'
fine_tune_checkpoint = ModelCheckpoint(fine_tune_checkpoint_path,
monitor='val_accuracy', verbose=1,
                                         save_best_only=True, mode='max')

print("\nFine-tuning CNN model...")
history_fine = model.fit(
    X_train, y_train,
    validation_split=0.2,
    epochs=5,
    batch_size=32,
    verbose=2,
    callbacks=[fine_tune_checkpoint]
)

# Load the best weights from fine-tuning
model.load_weights(fine_tune_checkpoint_path)

# 7. EVALUATE MODEL ON TEST SET
print("\nEvaluating model on test set:")
y_pred_proba = model.predict(X_test).flatten()
y_pred = (y_pred_proba > 0.5).astype(int)

print(classification_report(y_test, y_pred))
print("ROC-AUC Score:", roc_auc_score(y_test, y_pred_proba))

```

OUTPUT:

```

Extracting zip file...
Extracted Folders: ['Readme.txt', 'Forgery', 'Genuine']
Genuine examples: ['original_48_12.png', 'original_55_7.png', 'original_27_2.png']
Forgery examples: ['forgeries_27_23.png', 'forgeries_19_7.png', 'forgeries_11_5.png']

Loading and preprocessing dataset...
Loaded 2640 samples.

Splitting into train/test...

Training CNN model...
Epoch 1/5

Epoch 1: val_accuracy improved from -inf to 0.71868, saving model to best_model_initial.keras
53/53 - 361s - 7s/step - accuracy: 0.6294 - loss: 0.7639 - val_accuracy: 0.7187 - val_loss: 0.5397
Epoch 2/5

Epoch 2: val_accuracy improved from 0.71868 to 0.77778, saving model to best_model_initial.keras
53/53 - 366s - 7s/step - accuracy: 0.6951 - loss: 0.6292 - val_accuracy: 0.7778 - val_loss: 0.4692
Epoch 3/5

Epoch 3: val_accuracy improved from 0.77778 to 0.79669, saving model to best_model_initial.keras
53/53 - 394s - 7s/step - accuracy: 0.7016 - loss: 0.5912 - val_accuracy: 0.7967 - val_loss: 0.4406
Epoch 4/5

Epoch 4: val_accuracy improved from 0.79669 to 0.81087, saving model to best_model_initial.keras
53/53 - 348s - 7s/step - accuracy: 0.7063 - loss: 0.5806 - val_accuracy: 0.8109 - val_loss: 0.4164
Epoch 5/5

Epoch 5: val_accuracy did not improve from 0.81087
53/53 - 379s - 7s/step - accuracy: 0.7620 - loss: 0.4990 - val_accuracy: 0.8061 - val_loss: 0.4212

Fine-tuning CNN model...
Epoch 1/5

Epoch 1: val_accuracy improved from -inf to 0.70213, saving model to best_model_finetuned.keras
53/53 - 451s - 9s/step - accuracy: 0.7507 - loss: 0.5225 - val_accuracy: 0.7021 - val_loss: 0.6594
Epoch 2/5

Epoch 2: val_accuracy improved from 0.70213 to 0.80851, saving model to best_model_finetuned.keras
53/53 - 425s - 8s/step - accuracy: 0.8336 - loss: 0.3893 - val_accuracy: 0.8085 - val_loss: 0.4242
Epoch 3/5

Epoch 3: val_accuracy improved from 0.80851 to 0.88180, saving model to best_model_finetuned.keras
53/53 - 423s - 8s/step - accuracy: 0.8688 - loss: 0.3242 - val_accuracy: 0.8818 - val_loss: 0.2963
Epoch 4/5

```

```

Epoch 3: val_accuracy improved from 0.80851 to 0.88180, saving model to best_model_finetuned.keras
53/53 - 423s - 8s/step - accuracy: 0.8680 - loss: 0.3242 - val_accuracy: 0.8818 - val_loss: 0.2963
Epoch 4/5

Epoch 4: val_accuracy improved from 0.88180 to 0.90544, saving model to best_model_finetuned.keras
53/53 - 457s - 9s/step - accuracy: 0.9130 - loss: 0.2504 - val_accuracy: 0.9054 - val_loss: 0.2482
Epoch 5/5

```

```

Epoch 5: val_accuracy improved from 0.90544 to 0.91962, saving model to best_model_finetuned.keras
53/53 - 445s - 8s/step - accuracy: 0.9390 - loss: 0.1972 - val_accuracy: 0.9196 - val_loss: 0.2142

```

Evaluating model on test set:

17/17		94s 5s/step		
		precision	recall	f1-score
	0	0.92	0.85	0.88
	1	0.86	0.93	0.89
accuracy				0.89
macro avg		0.89	0.89	0.89
weighted avg		0.89	0.89	0.89

ROC-AUC Score: 0.9570850550964187

FGSM Attack Implementation CODE:

```
import tensorflow as tf
import numpy as np
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.applications.resnet50 import ResNet50,
preprocess_input
from sklearn.metrics import classification_report, roc_auc_score
import matplotlib.pyplot as plt

# 1. Build your model exactly as during training/fine-tuning
base_model = ResNet50(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
predictions = Dense(1, activation='sigmoid')(x)
model = Model(inputs=base_model.input, outputs=predictions)

# 2. Load best fine-tuned weights
model.load_weights('best_model_finetuned.keras')

# 3. Shape-correct FGSM batch implementation
def fgsm_attack_batch(model, images, labels, epsilon=0.01, batch_size=1):
    adv_images = []
    for start in range(0, len(images), batch_size):
        end = min(start + batch_size, len(images))
        img_batch = tf.convert_to_tensor(images[start:end])
        lbl_batch = tf.convert_to_tensor(labels[start:end],
dtype=tf.float32)
        lbl_batch = tf.reshape(lbl_batch, (-1, 1)) # Guarantee
(batch_size, 1)
        with tf.GradientTape() as tape:
            tape.watch(img_batch)
            pred = model(img_batch)
            pred = tf.reshape(pred, (-1, 1)) # Guarantee
(batch_size, 1)
            loss = tf.keras.losses.binary_crossentropy(lbl_batch, pred)
            grad = tape.gradient(loss, img_batch)
            adv = img_batch + epsilon * tf.sign(grad)
            adv = tf.clip_by_value(adv, -1.0, 1.0)
            adv_images.append(adv.numpy())
    return np.concatenate(adv_images, axis=0)

# 4. Test on a small batch for safety
NUM_SAMPLES = 20 # Start small for memory safety
x_test_small = x_test[:NUM_SAMPLES]
y_test_small = y_test[:NUM_SAMPLES]
epsilon_val = 0.01
```

```

X_test_adv = fgsm_attack_batch(model, X_test_small, y_test_small,
epsilon=epsilon_val, batch_size=1)

# 5. Evaluate
y_pred_adv_proba = model.predict(X_test_adv, batch_size=1).flatten()
y_pred_adv = (y_pred_adv_proba > 0.5).astype(int)
print("Evaluation on adversarial examples (FGSM, epsilon = %s):" % epsilon_val)
print(classification_report(y_test_small, y_pred_adv))
print("ROC-AUC Score on adversarial examples:",
roc_auc_score(y_test_small, y_pred_adv_proba))

# 6. Visualize
def visualize_adversarial(original, adversarial, n=5):
    plt.figure(figsize=(12, 4))
    for i in range(n):
        plt.subplot(2, n, i + 1)
        plt.imshow((original[i] + 1) / 2)
        plt.title("Original")
        plt.axis('off')
        plt.subplot(2, n, i + 1 + n)
        plt.imshow((adversarial[i] + 1) / 2)
        plt.title("Adversarial")
        plt.axis('off')
    plt.show()

visualize_adversarial(X_test_small, X_test_adv, n=min(5, NUM_SAMPLES))

# Now you can safely scale up NUM_SAMPLES or batch size if successful!

```

CODE OUTPUT:

```

20/20      6s 17ms/step
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control
    warn_prf_average, modifier, f"(metric.capitalize() is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control
    warn_prf_average, modifier, f"(metric.capitalize() is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control
    warn_prf_average, modifier, f"(metric.capitalize() is", len(result))
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..76.0305].
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..76.0305].
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..76.0305].
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..76.0305].
Evaluation on adversarial examples (FGSM, epsilon = 0.01):
    precision    recall   f1-score   support
          0       0.00     0.00     0.00      10
          1       0.50     0.67     0.67      10
    accuracy         0.50      20
   macro avg       0.25     0.50     0.33      20
weighted avg       0.25     0.50     0.33      20

```

ROC-AUC Score on adversarial examples: 0.4100000000000003

Adversarial Training CODE:

```

import tensorflow as tf
import numpy as np

```

```

# Use a small subset for demonstration (first 10 samples)
X_train_small = X_train[:5]
y_train_small = y_train[:5]

def generate_adversarial_examples(model, images, labels, epsilon=0.01):
    images = tf.convert_to_tensor(images)
    labels = tf.convert_to_tensor(labels, dtype=tf.float32)
    labels = tf.reshape(labels, (-1, 1))
    with tf.GradientTape() as tape:
        tape.watch(images)
        predictions = model(images)
        loss = tf.keras.losses.binary_crossentropy(labels, predictions)
        grads = tape.gradient(loss, images)
        adv_images = images + epsilon * tf.sign(grads)
        adv_images = tf.clip_by_value(adv_images, -1.0, 1.0)
    return adv_images.numpy()

# Compile the model if not compiled
model.compile(optimizer=tf.keras.optimizers.Adam(1e-4),
loss='binary_crossentropy', metrics=['accuracy'])

epochs = 1           # For demo; increase if desired
batch_size = 1        # Smallest, safest size
epsilon = 0.01        # Perturbation

for epoch in range(epochs):
    print(f"Adversarial Training Epoch {epoch+1}/{epochs}")
    # Shuffle only the small set
    idx = np.random.permutation(len(X_train_small))
    X_train_shuffled = X_train_small[idx]
    y_train_shuffled = y_train_small[idx]
    for start in range(0, len(X_train_small), batch_size):
        end = min(start + batch_size, len(X_train_small))
        clean_images = X_train_shuffled[start:end]
        labels = y_train_shuffled[start:end]
        adv_images = generate_adversarial_examples(model, clean_images,
labels, epsilon)
        X_combined = np.concatenate([clean_images, adv_images])
        y_combined = np.concatenate([labels, labels])
        model.train_on_batch(X_combined, y_combined)

# Save your robust (adversarially trained) model weights
model.save_weights('robust_model_adversarial_trained_demo.weights.h5')
print("Robust model weights (demo) saved as
'robust_model_adversarial_trained_demo.weights.h5'")

```

CODE OUTPUT:

```
⚡ Adversarial Training Epoch 1/1
Robust model weights (demo) saved as 'robust_model_adversarial_trained_demo.weights.h5'
```

Evaluation and Visualization OF THAT TRAINED MODEL CODE:

```
# 1. Rebuild the SAME model architecture as always
base_model = ResNet50(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
predictions = Dense(1, activation='sigmoid')(x)
robust_model = Model(inputs=base_model.input, outputs=predictions)

# 2. Load the robust model weights
robust_model.load_weights('robust_model_adversarial_trained_demo.weights.h5')

# 3. Compile for evaluation
robust_model.compile(optimizer=tf.keras.optimizers.Adam(1e-4),
loss='binary_crossentropy', metrics=['accuracy'])

# 4. Generate adversarial test examples using the new robust model
def fgsm_attack_batch(model, images, labels, epsilon=0.01, batch_size=1):
    adv_images = []
    for start in range(0, len(images), batch_size):
        end = min(start + batch_size, len(images))
        img_batch = tf.convert_to_tensor(images[start:end])
        lbl_batch = tf.convert_to_tensor(labels[start:end],
                                         dtype=tf.float32)
        lbl_batch = tf.reshape(lbl_batch, (-1, 1))
        with tf.GradientTape() as tape:
            tape.watch(img_batch)
            pred = model(img_batch)
            loss = tf.keras.losses.binary_crossentropy(lbl_batch, pred)
            grad = tape.gradient(loss, img_batch)
            adv = img_batch + 0.01 * tf.sign(grad)
            adv = tf.clip_by_value(adv, -1.0, 1.0)
            adv_images.append(adv.numpy())
    return np.concatenate(adv_images, axis=0)

NUM_SAMPLES = 10
X_test_small = X_test[:NUM_SAMPLES]
y_test_small = y_test[:NUM_SAMPLES]
X_test_adv = fgsm_attack_batch(robust_model, X_test_small, y_test_small,
epsilon=0.01, batch_size=1)

# 5. Evaluate on adversarial images
y_pred_adv_proba = robust_model.predict(X_test_adv,
batch_size=1).flatten()
```

```

y_pred_adv = (y_pred_adv_proba > 0.5).astype(int)
print("Robust model evaluation on FGSM adversarial examples:")
print(classification_report(y_test_small, y_pred_adv))
print("ROC-AUC:", roc_auc_score(y_test_small, y_pred_adv_proba))

# 6. (Optional) Visualize before/after as before
def visualize_adversarial(original, adversarial, n=5):
    import matplotlib.pyplot as plt
    plt.figure(figsize=(12, 4))
    for i in range(n):
        plt.subplot(2, n, i+1)
        plt.imshow((original[i] + 1) / 2)
        plt.title("Original")
        plt.axis('off')
        plt.subplot(2, n, i+1+n)
        plt.imshow((adversarial[i] + 1) / 2)
        plt.title("Adversarial")
        plt.axis('off')
    plt.show()
visualize_adversarial(X_test_small, X_test_adv, n=min(5, NUM_SAMPLES))

```

CODE OUTPUT:

```

10/10 4s 160ms/step
Robust model evaluation on FGSM adversarial examples:
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control
    _warn_prf(average, modifier, f"({metric.capitalize()}() is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control
    _warn_prf(average, modifier, f"({metric.capitalize()}() is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control
    _warn_prf(average, modifier, f"({metric.capitalize()}() is", len(result))
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..-76.0395].
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..-76.0395].
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..-76.0395].
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-61.34..-76.0395].
precision    recall    f1-score   support
          0       0.00     0.00      0.00      4
          1       0.60     1.00      0.75      6
   accuracy                           0.60      10
  macro avg       0.30     0.50      0.38      10
weighted avg       0.36     0.60      0.45      10
ROC-AUC: 0.5

```

	Original	Original	Original	Original	Original
Original					
Adversarial					

References

1. He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep Residual Learning for Image Recognition. *CVPR*. <https://doi.org/10.1109/CVPR.2016.90>
2. Goodfellow, I. J., Shlens, J., & Szegedy, C. (2015). Explaining and Harnessing Adversarial Examples. *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1412.6572>
3. Simonyan, K., & Zisserman, A. (2015). Very Deep Convolutional Networks for Large-Scale Image Recognition. *ICLR*. <https://arxiv.org/abs/1409.1556>
4. TensorFlow Documentation, Adversarial Example using FGSM. https://www.tensorflow.org/tutorials/generative/adversarial_fgsm
5. CEDAR Signature Dataset. Center of Excellence for Document Analysis and Recognition. <https://cedar.buffalo.edu/signature/>
6. Kaggle Signature Verification Dataset. <https://www.kaggle.com/datasets/tienen/handwritten-signature-verification>
7. Szegedy, C. et al. (2014). Intriguing properties of neural networks. <https://arxiv.org/abs/1312.6199>
8. Papernot, N., et al. (2016). Distillation as a Defense to Adversarial Perturbations against Deep Neural Networks. *IEEE Symposium on Security and Privacy*. <https://arxiv.org/abs/1511.04508>